

Beaver Works Summer Institute 2018

Unmanned Aerial Vehicle

Spencer Ng

Drone Architecture

- Two main control systems
- AeroFC responsible for just flight, very low level (PX4)
- Aero Compute Board is high level, using ROS as middleware

Robotics Operating System (ROS)

- Use `rostopic echo <topic>` to print from a topic
- MAVROS is for AeroFC, ROS is for Aero Compute Controller
- Nodes subscribe and publish to topics
 - Make decisions based on messages sent by other nodes

Localization Techniques

- Need to recognize where the drone is in the room or relative to frame of reference (origin)
- Includes velocity, pitch, roll
- Techniques include
 - Dead reckoning
 - * Using gyros, accelerometers, pots to determine how much a plant has moved
 - * Not good due to external inputs (e.g. wind)
 - GPS
 - * Gives a very good signal
 - * Precision only within meters
 - * Can't use indoors
 - Radio Frequency/Ultra Wideband
 - * Good for controlled environments, needs to be set up
 - Motion capture
 - * Fast and reliable
 - * Very popular for research due to high precision
 - * Only good for dedicated rooms and expensive
 - * Unrealistic
 - Visual Odometry
 - * How we navigate, with a visual map

- SLAM (Simultaneous Localization And Mapping)
 - * Creates a map and tries to localize at the same time
 - * Huge issue in robotics today, in terms of accuracy and performance
 - * IntelAero is capable, but is difficult to achieve well
 - * Active area of research

Reference Frames

- Need to understand quadrotor reference frame with respect to the world
- Either inertial or non-inertial (not accelerating, may be at constant nonzero velocity)
- NED or ENU is the world reference frame
 - $+z$ is down in NED (north east down)
 - $+z$ is up in ENU (east north up)
 - $+x$ is north and $+y$ is east in NED, vice versa in ENU
- Body frame (quadrotor)
 - $+x$ toward front
 - $+z$ down
 - $+y$ toward right "wing"
 - Origin is at center of mass
- Marker frame is attached to AR marker
 - Why tags are so useful
- Transformations can be made mathematically

Convolutions

- Moving average of intensity values in a grid produces edges
 - Smoothing function to suppress noise but makes it blurrier
- Image segmentation based on threshold
 - If above, make white; otherwise, change to black
- Shift invariance
- Linear filters
- To handle boundaries, duplicating parts of the image rather than leaving an edge produces a better result after running blurring kernel
- Kernels are matrices for filters
 - Sobels are for edge detection, starting from a particular direction
 - New pixel intensity is equal to "cross product" of kernel with adjacent pixels, equal to kernel size
- Linear regression - extract lines out of an image
 - Algebraic approach involves computing averages of the points to minimize squared error
 - Machine learning: guess and check from points (RANSAC technique)
 - * Two parameters: slope and intercept

- * Draw a line through two points at random, set "thresholds," determine how many "inliers" and outliers
 - Measure distance from every point to that line (either vertical/least squares or normal)
 - Distance is the cost to minimize
- * Rinse and repeat until inliers increase in number, until there's a good line
- * "Hopping around" the slope vs intercept vs squared error graph, for a set number of iterations
- * Convex optimization – not all problems have a definite curve to minimize error
 - Hard to change parameters and descend "cost hill"
- Operations
 - Adding
 - Blending
 - Erosion
 - Dilation
 - Rotation

Computer Vision

- AR tags in use
- A .bag file stores raw sensor data
- More than light and scene properties?
 - Optical illusions and human perception is important
- Edge detection
 - Use intensity rather than color differences
 - Makes us capable of recognizing objects
 - Horizon and vanishing points used for depth perception
 - Surface discontinuity – based on light, depth, or shadows
- Computers better than humans at parallel processing
- Goal of vision is to turn binary data into something recognizable
 - Many applications
- 3D graph of intensity as a grid of pixels
- Image is still a discrete signal, rather than continuous
- Intensity ranges from 0 to 255
- Color models are RGB, lab, or HSV
- An image can be a function from \mathbb{R}^2 to \mathbb{R}^M , where M is the number of colors
 - A position (x, y) is converted into (R, G, B)
- Images formed either perspective projection (has diagonals) or orthographic projection (only parallel lines)
 - Orthographic taken from very far away, then zooming in
- Image formation model for three dimensional world
 - Y and Z axes are meshed together in 2D representation

- X is preserved, in the same orientation
- Edge classification
 - Figure segmentation
 - Occlusion edges
 - * Only the foreground
 - Contact edges – touching the ground
 - * Y is 0 because the object is touching the ground
 - * Z value is still unknown
- Accidental image – distorts perception of axes, loses info
- Differential geometry - using $(N - 1)$ -D images to visualize N -D objects
- Systems and filters
 - Goal is to extract feature (edges, blobs) from images
 - Denoising, in-painting, etc
 - Linear filter in 1D
 - * A set of n values in an array g
 - * Another signal (sort of like modulating in AM) is introduced to filter/transform that signal
 - * Used to create an overall transformative function, $f(n) = H(g(n), s(x, y))$
 - Translation invariant filter – do same operation at every point in the image to filter it
 - * If input image is filtered by m samples, output also translated by m samples
 - Convolution
 - * Key operation for machine learning/embedded (local) vision
 - * Sliding/multiplying a set of values across another set of values to produce a transformation
 - * Same as cross correlation
- Characterizing edges – place of rapid change in image intensity
 - First derivative of intensity vs. Position graph (aka gradient profile) shows spikes/peaks
 - Noise causes derivative/gradient profile to be unreadable, as there are many edges
 - Finite differences is susceptible to noise
 - Smooth/filter by combining the input signal and a bell curve (Gaussian distribution) $f(x) = s(b(x))$
 - To detect change in intensity, just take the derivative and use convolution after smoothing
 - Power rule: $\frac{d}{dx}x^n = n \cdot x^{n-1}$
- Deep learning
 - Convolution Neural Networks (CNN) - collections of basic mathematical operations
 - Training begins with discriminator and generator, which learn from each other
 - Called "deep" because transformations keep on going further down
- Recurrent Neural Network (RNN)
 - Keep on transforming the image; feed output back into input

Color Segmentation

- RGB color space – visualized as 3D graph, values 0-255 for each channel
- HSL/HSV - polar coordinates, visualized as cylinder
 - Saturation is radius
 - Lightness/value is height
 - Hue is degree/direction
- Color thresholding – cutoff at particular value
- Color masking

States

- State vector contains 12 variables for complete and unique description of quadrotor
 - $\dot{\vec{\zeta}} = f(\vec{\zeta}, \vec{u})$
 - x, y, z for position (could be very noisy)
 - v_x, v_y, v_z for velocity (often calculated or estimated)
 - ϕ, θ, ψ for roll, pitch, and yaw
 - p, q, r for pitch, roll, and yaw "velocities"

$$\vec{\zeta} = \begin{bmatrix} x \\ y \\ z \\ v_x \\ v_y \\ v_z \\ \phi \\ \theta \\ \psi \\ p \\ q \\ r \end{bmatrix}$$

- \vec{u} contains angular velocities for four rotors ($\omega_1, \omega_2, \omega_3, \omega_4$)
 - Derive acceleration/force using $F = m \cdot \vec{a}$ and $\tau = I \cdot \dot{\omega}$

$$\vec{u} = \begin{bmatrix} \omega_1 \\ \omega_2 \\ \omega_3 \\ \omega_4 \end{bmatrix} \rightarrow \begin{bmatrix} \omega_1^2 \\ \omega_2^2 \\ \omega_3^2 \\ \omega_4^2 \end{bmatrix} \rightarrow \begin{bmatrix} f \\ \tau_1 \\ \tau_2 \\ \tau_3 \end{bmatrix}$$

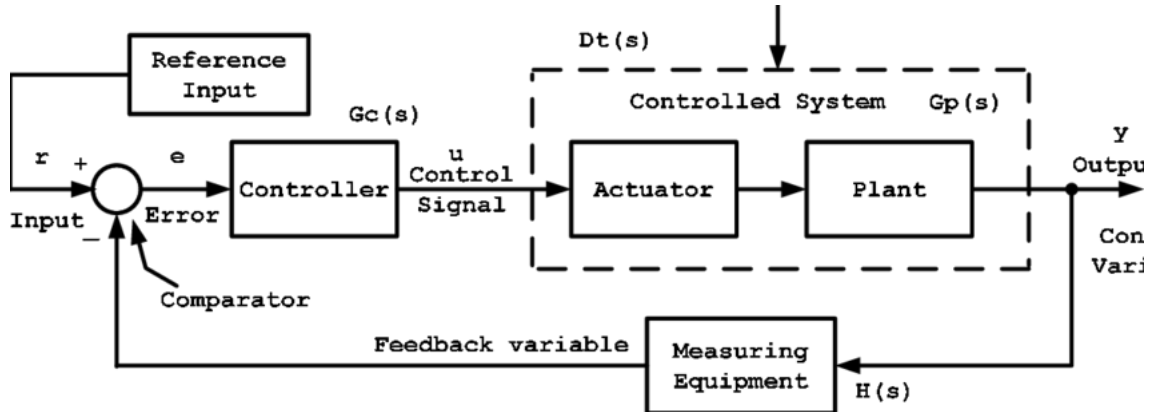
- Above is the low level control, where f is the sum of all forces on four rotors, τ_n is the torque on the center of mass on the quadrotor, in all three directions
- Rightmost vector is what's really used
- Physics is mostly handled by the flight controller, which we just send velocity commands to via MAVROS

$$\begin{aligned}\dot{\vec{x}} &= \frac{d^w \vec{x}_B}{dt} \\ \ddot{x} &= mg\vec{z}_w - u_1 \vec{z}_b \\ \dot{R}_{BW} &= R_{BW} \vec{\Omega}_{BW} \\ \dot{\vec{\Omega}}_{BW} &= \begin{bmatrix} u_2 \\ u_3 \\ u_4 \end{bmatrix} - \vec{\Omega}_{BW} \times J_B \vec{\Omega}_{BW}\end{aligned}$$

Controls & Estimation

- Three aspects toward autonomous control
 - State estimation
 - Planning
 - Controls
- Use the same closed control loop with controller (feeds inputs), UAV (provides outputs from sensors), and observer (state estimation)
- Problem is that sensors have noise
 - Inertia measurement unit (IMU) - measures acceleration and velocity, but prone to drift over time
 - * Error also accumulates over time
 - GPS - current position, but noisy and often has large error due to lack of precision
- Need to develop mathematical model to figure out the state
- Simple state estimator - average values of particular variable derived from three sensors, depending on your confidence in them
- Mathematical model should be derived to calculate uncertainty also
 - Only an approximation
 - Subject to uncertainties
- If there's no uncertainty, then $y = \hat{y}$
- State observer
 - Help observe current state of rotor and brings actual results toward ideal mathematical model
- Kalman Filter
 - Optimal estimation algorithm to detect location and speed in noisy measurements
 - First used in 1960 for Apollo project
 - Used when indirect measurement is available, but not direct measurement
 - Combines measurements and predictions
 - Two step process:
 - * Prediction - system model used to calculate prior state
 - * Update - use measurements to calculate and refine the prediction
 - * Final output is a weighted process, depending on your confidence of model ($x = \alpha_1 x_1 + \alpha_2 x_2$)
 - Can only be used for linear systems
- Nonlinear state estimators needed for quadrotors
 - Use an extended Kalman Filter - one for each linear variable

Feedback Control



- Terminology
 - **Plant** - system to be controlled, the "physics" of the problem (e.g. UAV)
 - **Sensor** - measures quantities
 - **Actuator** - controls the plant (e.g. motor)
 - **Controller** - processors sensor data to drive actuator
 - **Control law** - algorithm used by control processor to compute actuator signal (math)
- Example: hot plate temperature
 - Equation $T = T_{\text{amb}} + \alpha(P + D)$ is used for output, where P is input power, D is external disturbance, α is in $^{\circ}\text{F}/\text{W}$
 - Using algebra, manipulate to solve for P , given a T_{desired} - $P = \frac{T_{\text{des}} - T_{\text{amb}}}{\alpha}$, assuming D is 0 W.
 - Error is $e = T - T_{\text{des}}$, goal is to have this be zero
 - In open loop control, have to manually adjust, with trial and error
 - Closed loop control equation: $P = \frac{T_{\text{des}} - T_{\text{amb}}}{\alpha} - k(T - T_{\text{des}})$
 - * Increase power if too cold, lower power if too warm
 - * Feed the error back into the system
 - * The k term can be optimized with trial and error
 - In the end,

$$\text{closed-loop error} = \frac{\text{open-loop error}}{1 + \alpha k}$$
 - Error is significantly less in closed-loop control when getting to a steady state
- Slide camera example
 - A camera moves along a fixed track with velocity \vec{v} and camera pitch angle θ , and its target position x_{des} is at an angle γ with respect to θ
 - $\dot{x} = \vec{v}$ and $\dot{\theta} = K \sin \theta$
 - Can detect $\gamma = -(\tan^{-1}(\frac{x - x_{\text{des}}}{h}) + \theta)$ based on sensors in the camera frame, where x is the current 2D position of the drone along the ground and h is its height
 - Letting $u = \theta$ (controller to plant value), goal is to find $\theta(t)$ to control the drone that makes $\gamma \rightarrow \gamma_{\text{des}}$
 - * $\gamma_{\text{des}} = 0$ is the general goal
 - In a simplified version, $\theta = 0$ is effectively ignored, and goal is to drive $\dot{x} = \vec{v}$ and find $\vec{v}(t)$
 - * $\gamma = \tan^{-1}(\frac{x - x_{\text{des}}}{h})$, the angle between x_{des} and x with respect to the drone
 - * Find an optimized $u = \vec{v}$ that drives $\gamma \rightarrow \gamma_{\text{des}}$
 - * For a velocity command, $\vec{v} = K_p \gamma_{\text{err}}$
 - * For an acceleration command, $\vec{a} = K_p \cdot \frac{d}{dt} \gamma_{\text{err}}(t)$
 - * Integration of $\gamma_{\text{err}}(t)$ can be used to fix induced noise when sending velocity commands (e.g. due to wind)

PID Control

$$u(t) = \underbrace{K_P \cdot e(t)}_{\text{current}} + \underbrace{K_D \frac{d}{dt} e(t)}_{\text{future}} + \underbrace{K_I \int_0^t e(t)}_{\text{past}}$$

- Choosing gains
 1. Increase K_P until system starts to oscillate around the target, with other gains at 0
 2. Increase K_D to damp motion/oscillation
 3. Increase K_I to get rid of steady state error
- P is almost universal in all controllers
 - May cause overshoot and oscillation, however
- I resolves steady-state error
 - Risks causing saturation
- D predicts future to avoid overshoot
 - Noise can exaggerate derivative commands

Optical Flow

- Motion field - 2D projection of 3D motion
- Optical flow field - 2D, apparent motion; requires reflection
- Despite 6 degrees of freedom, can only capture a 2D field
- Example: matte sphere rotating has motion field, but not an optical flow field
- Process of visualization
 - Two frames are given as input
 - Differences are computed and visualized
- Applications
 - Video stabilization
 - Denoising
 - Super resolution
- Scanning locally (small window) for changes is more effective
- Goal is to find differences between two frames
- Need to find corners and determine if they are same
 - "Flat" region- no change in all directions
 - "Edge" - no change along edge direction
 - "Corner" - change in all directions (in uniform color shape)
 - Differences due to aliasing in images
- Windows - spacial and temporal
 - u is translation in x direction, v is "wiggle" in y
 - Not moving window around, but see how contents will change over time

$$E(u, v) = \sum_{x, y} \underbrace{w(x, y)}_{\text{window function}} \underbrace{[I(x + u, y + v) - I(x, y)]}_{\text{shifted intensity} - \text{intensity}}^2$$

- Change in appearance for the shift, as a function of pixel intensities
- Trial and error required to determine appropriate u and v
- Window function is a weighted sum of the pixels in the window
- u and v are $(0, 0)$ if the image is the same color throughout
- Small but nonzero shifts in u and v are interesting to find out, key to edge detection
- Taylor series - approximation of a function by taking the derivative throughout its input
- Product rule: $\frac{d}{dx} f(x)g(x) = f'(x)g(x) + f(x)g'(x)$

- For an image corner,

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} \approx \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

- Eigenvalues help define shape of an ellipse, in terms of stretch/compression factors
 - λ_1 is the direction of the fastest change, independent of λ_2
- Harris Detector Algorithm
 - Has translation and rotation invariance
 - No scalar invariance - zooming in on a corner reveals multiple "edges"
 - Detects glare and lighting changes easily, yet still identify object
- "Phase diagram" based on actual phase gives probabilistic, fast realtime estimation of object
- Optical flow has many methods
 - Pixel difference for every pixel
 - Corner detection
- Only is apparent motion of object
 - Problem occurs when motion vectors of individual pixels are in different directions
 - The pixel correspondence problem - how to estimate pixel motion
- Assumptions made in computing optical flow
 - Brightness/intensity of pixel is constant between frames (brightness constancy)
 - * Problem when lighting changes
 - * $I(x, y, t) = I(x + u, y + v, t + 1)$, once we know u and v
 - Solving equation with Taylor series, we get $I(x + u, y + v) = I(x, y, t + 1) + \frac{\partial I(t)}{\partial x} u + \frac{\partial I(t)}{\partial y} v$, but only applicable for small motions
 - Derivation from smoothness is Gaussian
- Aperture problem - if the window is too small, apparent motion direction may not be the actual direction
 - Solution is to pretend adjacent pixels have the same (u, v) in a window with n pixels

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ \vdots & \vdots \\ I_x(p_n) & I_y(p_n) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} I_{t_x}(p_1) & I_{t_y}(p_1) \\ \vdots & \vdots \\ I_{t_x}(p_n) & I_{t_y}(p_n) \end{bmatrix}$$

- Spatial smoothness - detect a whole surface in space, based on edges